

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns# data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list the files in the directory

import os

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session
```

```
In [2]: df = pd.read_csv('WA_Fn-UseC_-Telco-Customer-Churn.csv')
```

Understanding the Data

```
In [3]: df.head()
```

Out[3]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service
1	5575-GNVDE	Male	0	No	No	34	Yes	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service
4	9237-HQITU	Female	0	No	No	2	Yes	No

5 rows × 21 columns

```
In [4]: df.columns
```

Out[4]: Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'], dtype='object')

```
In [5]: df.drop("customerID", axis = 1, inplace = True)
```

```
In [6]: df.tenure.mean()
```

```
Out[6]: 32.37114865824223
```

```
In [7]: df.Contract.value_counts()
```

```
Out[7]: Month-to-month    3875
Two year                1695
One year                1473
Name: Contract, dtype: int64
```

```
In [8]: churn_rate = df.Churn[df.Churn == "Yes"].count() / df.shape[0]
print(churn_rate)
```

```
0.2653698707936959
```

```
In [9]: churned = df.loc[df.Churn == "Yes"]
```

```
In [10]: churned["Contract"].value_counts() #long contracts leads to lesser churn rate
```

```
Out[10]: Month-to-month    1655
One year                166
Two year                48
Name: Contract, dtype: int64
```

```
In [11]: churned["tenure"].value_counts() #similarly over here too
```

```
Out[11]: 1      380
2      123
3       94
4       83
5       64
...
60       6
72       6
62       5
64       4
63       4
Name: tenure, Length: 72, dtype: int64
```


```
In [12]: df.Churn.value_counts() #1 represents Churned
```

```
Out[12]: No      5174
Yes      1869
Name: Churn, dtype: int64
```

```
In [13]: CTR = pd.crosstab(index = df.Contract, columns = df.Churn)
```

```
In [14]: CTR.plot.barh()
```

```
Out[14]: <Axes: ylabel='Contract'>
```



```
In [15]: CTR
```

```
Out[15]:
```

	Churn	No	Yes
Contract			
Month-to-month		2220	1655
One year		1307	166
Two year		1647	48

```
In [16]: df.isnull().sum()
```

```
Out[16]: gender                0
SeniorCitizen                0
Partner                      0
Dependents                   0
tenure                       0
PhoneService                 0
MultipleLines                0
InternetService              0
OnlineSecurity               0
OnlineBackup                 0
DeviceProtection             0
TechSupport                  0
StreamingTV                  0
StreamingMovies              0
Contract                     0
PaperlessBilling             0
PaymentMethod                0
MonthlyCharges               0
TotalCharges                 0
Churn                        0
dtype: int64
```

```
In [17]: df.info() #Total charges is object data type, must be some white space since no
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                7043 non-null  object
1   SeniorCitizen         7043 non-null  int64
2   Partner               7043 non-null  object
3   Dependents            7043 non-null  object
4   tenure                7043 non-null  int64
5   PhoneService          7043 non-null  object
6   MultipleLines         7043 non-null  object
7   InternetService       7043 non-null  object
8   OnlineSecurity        7043 non-null  object
9   OnlineBackup          7043 non-null  object
10  DeviceProtection      7043 non-null  object
11  TechSupport           7043 non-null  object
12  StreamingTV           7043 non-null  object
13  StreamingMovies       7043 non-null  object
14  Contract              7043 non-null  object
15  PaperlessBilling      7043 non-null  object
16  PaymentMethod         7043 non-null  object
17  MonthlyCharges        7043 non-null  float64
18  TotalCharges          7043 non-null  object
19  Churn                 7043 non-null  object
dtypes: float64(1), int64(2), object(17)
memory usage: 1.1+ MB
```

```
In [18]: df.TotalCharges = pd.to_numeric(df.TotalCharges, errors = "coerce")
```

```
In [19]: df.TotalCharges #dtype is float now
```

```
Out[19]: 0      29.85
1    1889.50
2     108.15
3    1840.75
4     151.65
...
7038   1990.50
7039   7362.90
7040    346.45
7041    306.60
7042   6844.50
Name: TotalCharges, Length: 7043, dtype: float64
```

```
In [20]: df.isnull().sum() #now we have 11 missing values from the same column
```

```
Out[20]: gender      0
SeniorCitizen      0
Partner            0
Dependents         0
tenure             0
PhoneService       0
MultipleLines      0
InternetService    0
OnlineSecurity     0
OnlineBackup       0
DeviceProtection   0
TechSupport        0
StreamingTV        0
StreamingMovies    0
Contract           0
PaperlessBilling   0
PaymentMethod      0
MonthlyCharges     0
TotalCharges      11
Churn              0
dtype: int64
```

```
In [21]: mask = df.TotalCharges.isna()
missing = df[mask] #subsetting the missing rows
```

In [22]: missing

Out[22]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
488	Female	0	Yes	Yes	0	No	No phone service	
753	Male	0	No	Yes	0	Yes	No	
936	Female	0	Yes	Yes	0	Yes	No	
1082	Male	0	Yes	Yes	0	Yes	Yes	
1340	Female	0	Yes	Yes	0	No	No phone service	
3331	Male	0	Yes	Yes	0	Yes	No	
3826	Male	0	Yes	Yes	0	Yes	Yes	
4380	Female	0	Yes	Yes	0	Yes	No	
5218	Male	0	Yes	Yes	0	Yes	No	
6670	Female	0	Yes	Yes	0	Yes	Yes	
6754	Male	0	No	Yes	0	Yes	Yes	

In [23]: df.iloc[753,19] *#inspecting the missing values individually, white space it is*

Out[23]: 'No'

In [24]: *"""instead of using MICE or replacing it with median, logical replacement will
The tenure column is assumed to be number of months. Since the tenure is 0 for
values, thus 0 seems to be the most likely replacement"""*

Out[24]: 'instead of using MICE or replacing it with median, logical replacement will
be better.\nThe tenure column is assumed to be number of months. Since the te
nure is 0 for all missing\nvalues, thus 0 seems to be the most likely replace
ment'

In [25]: df.TotalCharges.fillna(0, inplace=True)

In [26]: df.TotalCharges.dtype

Out[26]: dtype('float64')

```
In [27]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                7043 non-null   object
1   SeniorCitizen         7043 non-null   int64
2   Partner               7043 non-null   object
3   Dependents            7043 non-null   object
4   tenure                7043 non-null   int64
5   PhoneService          7043 non-null   object
6   MultipleLines         7043 non-null   object
7   InternetService       7043 non-null   object
8   OnlineSecurity        7043 non-null   object
9   OnlineBackup          7043 non-null   object
10  DeviceProtection      7043 non-null   object
11  TechSupport           7043 non-null   object
12  StreamingTV           7043 non-null   object
13  StreamingMovies       7043 non-null   object
14  Contract              7043 non-null   object
15  PaperlessBilling      7043 non-null   object
16  PaymentMethod         7043 non-null   object
17  MonthlyCharges        7043 non-null   float64
18  TotalCharges          7043 non-null   float64
19  Churn                 7043 non-null   object
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB
```

```
In [28]: df.PaymentMethod.unique()
```

```
Out[28]: array(['Electronic check', 'Mailed check', 'Bank transfer (automatic)',
               'Credit card (automatic)'], dtype=object)
```

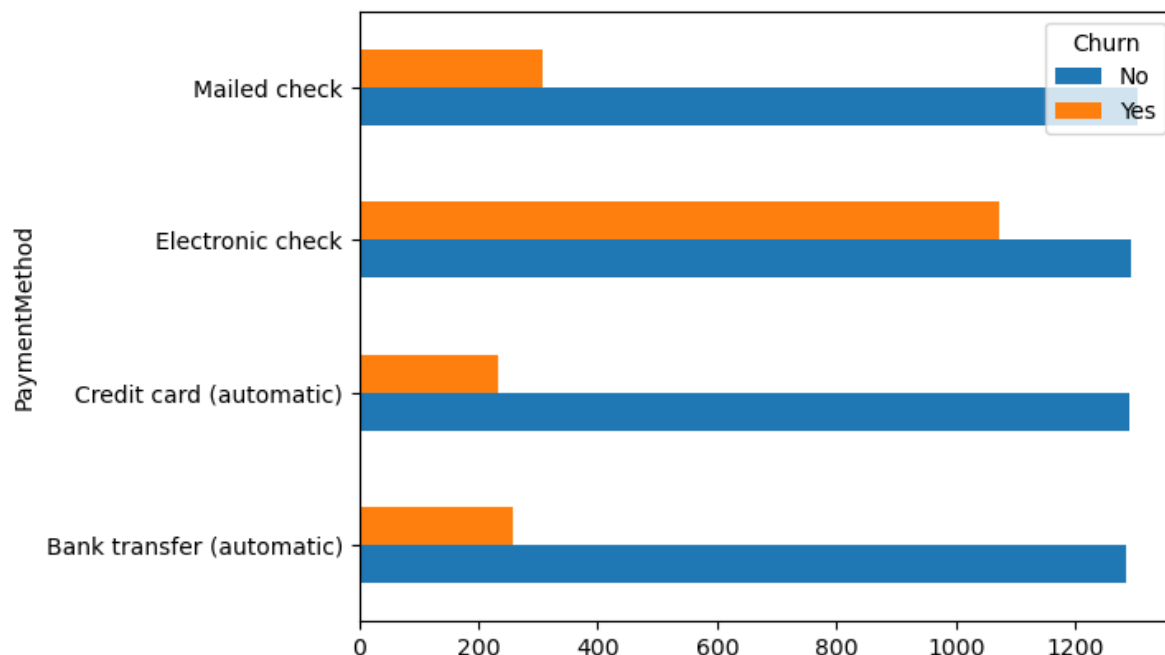
```
In [29]: df.PaymentMethod.value_counts()
```

```
Out[29]: Electronic check                2365
Mailed check                            1612
Bank transfer (automatic)               1544
Credit card (automatic)                1522
Name: PaymentMethod, dtype: int64
```

```
In [30]: CPR = pd.crosstab(index = df.PaymentMethod, columns= df.Churn)
```

```
In [31]: CPR.plot.barh()
```

```
Out[31]: <Axes: ylabel='PaymentMethod'>
```



```
In [32]: df.StreamingTV.value_counts()
```

```
Out[32]: No                2810  
Yes                2707  
No internet service    1526  
Name: StreamingTV, dtype: int64
```

```
In [33]: numeric_cols = df.select_dtypes(["float64", "int64"])
```


In [34]: numeric_cols

Out[34]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
0	0	1	29.85	29.85
1	0	34	56.95	1889.50
2	0	2	53.85	108.15
3	0	45	42.30	1840.75
4	0	2	70.70	151.65
...
7038	0	24	84.80	1990.50
7039	0	72	103.20	7362.90
7040	0	11	29.60	346.45
7041	1	4	74.40	306.60
7042	0	66	105.65	6844.50

7043 rows × 4 columns

Preprocessing and Model Selection

```
In [35]: from sklearn import preprocessing
from sklearn.impute import SimpleImputer
from sklearn.linear_model import SGDClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import f1_score
from sklearn.metrics import precision_recall_curve
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn import svm
from sklearn.inspection import DecisionBoundaryDisplay
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.compose import ColumnTransformer
from sklearn.metrics import PrecisionRecallDisplay
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.model_selection import RandomizedSearchCV
```

```
In [36]: unique_counts = df.select_dtypes("O").nunique()
binary_columns = unique_counts[unique_counts == 2].index.drop("Churn").tolist()
categorical_columns = unique_counts[unique_counts > 2].index.tolist()
target_column = "Churn"
```

```
In [37]: X = df.drop(target_column, axis=1)
y = df[target_column]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, stratify=y, random_state=42
)
X_train.head(5)
```

Out[37]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService
5557	Female	0	No	No	5	Yes	No	Fiber
2270	Female	1	No	No	3	Yes	No	Fiber
6930	Female	0	Yes	No	3	Yes	Yes	Fiber
2257	Female	0	No	No	60	Yes	Yes	Fiber
898	Female	0	No	No	12	Yes	No	Fiber

```
In [38]: transformer = ColumnTransformer(  
    [  
        ("scaler", StandardScaler(), ["MonthlyCharges", "TotalCharges", "tenure"]),  
        ("binary_encoder", preprocessing.OrdinalEncoder(), binary_columns),  
        ("ohe", preprocessing.OneHotEncoder(drop="first"), categorical_columns),  
    ],  
    remainder="passthrough",  
)  
  
transformer.fit(X_train)  
columns = transformer.get_feature_names_out()  
columns = list(map(lambda x: str(x).split("__")[-1], columns))  
  
X_train = pd.DataFrame(transformer.transform(X_train), columns=columns)  
X_test = pd.DataFrame(transformer.transform(X_test), columns=columns)
```

```
In [39]: label_encoder = preprocessing.LabelEncoder()  
label_encoder.fit(y_train)  
y_train = label_encoder.transform(y_train)  
y_test = label_encoder.transform(y_test)
```

Support Vector Classifier

```
In [40]: param_grid = {'C': [0.1, 1, 10, 100, 1000],  
    'gamma': [1, 0.1, 0.01, 0.001, 0.0001],  
    'kernel': ['rbf']}
```

```
In [41]: SVC_grid = GridSearchCV(svm.SVC(), param_grid, refit = True, verbose = 6, cv =  
SVC_grid.fit(X_train, y_train)
```

```
Fitting 3 folds for each of 25 candidates, totalling 75 fits
[CV 1/3] END .....C=0.1, gamma=1, kernel=rbf;; score=0.754 total time= 1
0.9s
[CV 2/3] END .....C=0.1, gamma=1, kernel=rbf;; score=0.753 total time= 1
0.8s
[CV 3/3] END .....C=0.1, gamma=1, kernel=rbf;; score=0.744 total time= 1
0.6s
[CV 1/3] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.806 total time=
6.2s
[CV 2/3] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.796 total time=
5.6s
[CV 3/3] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.780 total time=
6.2s
[CV 1/3] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.735 total time=
6.1s
[CV 2/3] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.738 total time=
6.7s
[CV 3/3] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.768 total time=
6.5s
[CV 1/3] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.735 total time=
4.9s
[CV 2/3] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.735 total time=
5.5s
[CV 3/3] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.735 total time=
5.5s
[CV 1/3] END ...C=0.1, gamma=0.0001, kernel=rbf;; score=0.735 total time=
5.5s
[CV 2/3] END ...C=0.1, gamma=0.0001, kernel=rbf;; score=0.735 total time=
4.8s
[CV 3/3] END ...C=0.1, gamma=0.0001, kernel=rbf;; score=0.735 total time=
5.1s
[CV 1/3] END .....C=1, gamma=1, kernel=rbf;; score=0.790 total time=
9.3s
[CV 2/3] END .....C=1, gamma=1, kernel=rbf;; score=0.787 total time= 1
1.7s
[CV 3/3] END .....C=1, gamma=1, kernel=rbf;; score=0.772 total time=
9.3s
[CV 1/3] END .....C=1, gamma=0.1, kernel=rbf;; score=0.819 total time=
5.6s
[CV 2/3] END .....C=1, gamma=0.1, kernel=rbf;; score=0.804 total time=
5.4s
[CV 3/3] END .....C=1, gamma=0.1, kernel=rbf;; score=0.786 total time=
5.7s
[CV 1/3] END .....C=1, gamma=0.01, kernel=rbf;; score=0.814 total time=
6.5s
[CV 2/3] END .....C=1, gamma=0.01, kernel=rbf;; score=0.802 total time=
6.0s
[CV 3/3] END .....C=1, gamma=0.01, kernel=rbf;; score=0.783 total time=
6.5s
[CV 1/3] END .....C=1, gamma=0.001, kernel=rbf;; score=0.735 total time=
6.4s
[CV 2/3] END .....C=1, gamma=0.001, kernel=rbf;; score=0.739 total time=
5.6s
[CV 3/3] END .....C=1, gamma=0.001, kernel=rbf;; score=0.772 total time=
5.8s
[CV 1/3] END .....C=1, gamma=0.0001, kernel=rbf;; score=0.735 total time=
5.7s
```


```
[CV 2/3] END .....C=1, gamma=0.0001, kernel=rbf;, score=0.735 total time=
5.6s
[CV 3/3] END .....C=1, gamma=0.0001, kernel=rbf;, score=0.735 total time=
5.8s
[CV 1/3] END .....C=10, gamma=1, kernel=rbf;, score=0.768 total time=
7.4s
[CV 2/3] END .....C=10, gamma=1, kernel=rbf;, score=0.771 total time=
9.0s
[CV 3/3] END .....C=10, gamma=1, kernel=rbf;, score=0.761 total time= 1
0.0s
[CV 1/3] END .....C=10, gamma=0.1, kernel=rbf;, score=0.804 total time=
7.4s
[CV 2/3] END .....C=10, gamma=0.1, kernel=rbf;, score=0.796 total time=
6.0s
[CV 3/3] END .....C=10, gamma=0.1, kernel=rbf;, score=0.780 total time=
6.7s
[CV 1/3] END .....C=10, gamma=0.01, kernel=rbf;, score=0.815 total time=
5.5s
[CV 2/3] END .....C=10, gamma=0.01, kernel=rbf;, score=0.795 total time=
5.0s
[CV 3/3] END .....C=10, gamma=0.01, kernel=rbf;, score=0.781 total time=
5.3s
[CV 1/3] END .....C=10, gamma=0.001, kernel=rbf;, score=0.815 total time=
5.8s
[CV 2/3] END .....C=10, gamma=0.001, kernel=rbf;, score=0.805 total time=
6.6s
[CV 3/3] END .....C=10, gamma=0.001, kernel=rbf;, score=0.785 total time=
6.4s
[CV 1/3] END ....C=10, gamma=0.0001, kernel=rbf;, score=0.735 total time=
6.1s
[CV 2/3] END ....C=10, gamma=0.0001, kernel=rbf;, score=0.739 total time=
8.5s
[CV 3/3] END ....C=10, gamma=0.0001, kernel=rbf;, score=0.772 total time=
7.3s
[CV 1/3] END .....C=100, gamma=1, kernel=rbf;, score=0.765 total time= 1
6.6s
[CV 2/3] END .....C=100, gamma=1, kernel=rbf;, score=0.767 total time= 1
3.5s
[CV 3/3] END .....C=100, gamma=1, kernel=rbf;, score=0.758 total time= 1
2.8s
[CV 1/3] END .....C=100, gamma=0.1, kernel=rbf;, score=0.760 total time= 1
4.2s
[CV 2/3] END .....C=100, gamma=0.1, kernel=rbf;, score=0.743 total time= 1
6.2s
[CV 3/3] END .....C=100, gamma=0.1, kernel=rbf;, score=0.744 total time= 1
5.8s
[CV 1/3] END .....C=100, gamma=0.01, kernel=rbf;, score=0.823 total time= 1
0.9s
[CV 2/3] END .....C=100, gamma=0.01, kernel=rbf;, score=0.806 total time=
8.9s
[CV 3/3] END .....C=100, gamma=0.01, kernel=rbf;, score=0.783 total time=
8.9s
[CV 1/3] END ....C=100, gamma=0.001, kernel=rbf;, score=0.819 total time=
9.4s
[CV 2/3] END ....C=100, gamma=0.001, kernel=rbf;, score=0.800 total time=
5.4s
[CV 3/3] END ....C=100, gamma=0.001, kernel=rbf;, score=0.786 total time=
```

```

5.4s
[CV 1/3] END ...C=100, gamma=0.0001, kernel=rbf;, score=0.816 total time=
5.9s
[CV 2/3] END ...C=100, gamma=0.0001, kernel=rbf;, score=0.804 total time=
6.4s
[CV 3/3] END ...C=100, gamma=0.0001, kernel=rbf;, score=0.783 total time=
7.4s
[CV 1/3] END .....C=1000, gamma=1, kernel=rbf;, score=0.753 total time= 2
0.6s
[CV 2/3] END .....C=1000, gamma=1, kernel=rbf;, score=0.758 total time= 1
7.1s
[CV 3/3] END .....C=1000, gamma=1, kernel=rbf;, score=0.757 total time= 1
7.9s
[CV 1/3] END .....C=1000, gamma=0.1, kernel=rbf;, score=0.731 total time= 2
3.7s
[CV 2/3] END .....C=1000, gamma=0.1, kernel=rbf;, score=0.724 total time= 2
4.5s
[CV 3/3] END .....C=1000, gamma=0.1, kernel=rbf;, score=0.715 total time= 1
9.9s
[CV 1/3] END ....C=1000, gamma=0.01, kernel=rbf;, score=0.813 total time= 2
1.3s
[CV 2/3] END ....C=1000, gamma=0.01, kernel=rbf;, score=0.797 total time= 1
3.5s
[CV 3/3] END ....C=1000, gamma=0.01, kernel=rbf;, score=0.772 total time= 1
3.3s
[CV 1/3] END ...C=1000, gamma=0.001, kernel=rbf;, score=0.818 total time=
6.2s
[CV 2/3] END ...C=1000, gamma=0.001, kernel=rbf;, score=0.795 total time=
5.4s
[CV 3/3] END ...C=1000, gamma=0.001, kernel=rbf;, score=0.782 total time=
6.2s
[CV 1/3] END ..C=1000, gamma=0.0001, kernel=rbf;, score=0.816 total time=
5.4s
[CV 2/3] END ..C=1000, gamma=0.0001, kernel=rbf;, score=0.804 total time=
5.3s
[CV 3/3] END ..C=1000, gamma=0.0001, kernel=rbf;, score=0.781 total time=
4.8s

```


Out[41]:



In [42]: `y_pred_SVC = SVC_grid.predict(X_test)`

In [43]: `print(accuracy_score(y_test, y_pred_SVC))` *#rbf kernel gives the best accuracy*
 0.7998106956933271

```
In [44]: display = PrecisionRecallDisplay.from_estimator(
          SVC_grid, X_test, y_test, name="RbfSVC"
        )
        _ = display.ax_.set_title("2-class Precision-Recall curve") #Average precision
```



```
In [45]: print(classification_report(y_test, y_pred_SVC)) #model precision and recall na
```

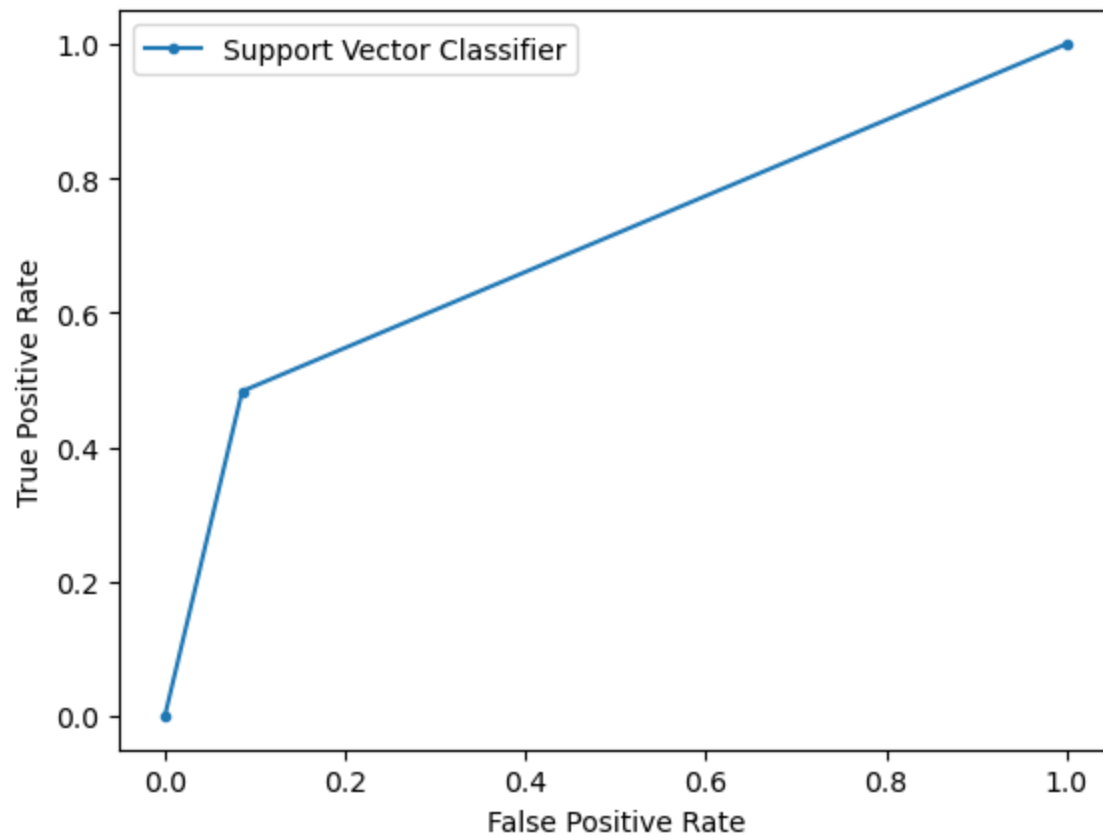
	precision	recall	f1-score	support
0	0.83	0.91	0.87	1552
1	0.67	0.48	0.56	561
accuracy			0.80	2113
macro avg	0.75	0.70	0.72	2113
weighted avg	0.79	0.80	0.79	2113

```
In [46]: roc_auc_score(y_test, y_pred_SVC)
```

```
Out[46]: 0.6986850386827647
```




```
In [47]: lr_fpr_SVC, lr_tpr_SVC, _ = roc_curve(y_test, y_pred_SVC)
# plot the roc curve for the model
plt.plot(lr_fpr_SVC, lr_tpr_SVC, marker='.', label='Support Vector Classifier')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
# show the legend
plt.legend()
# show the plot
plt.show()
```



```
In [48]: confusion_matrix_SVC = confusion_matrix(y_test, y_pred_SVC)

cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix_SVC)

cm_display.plot()
plt.show()
```



Decision Tree Classifier

```
In [49]: params = {'max_leaf_nodes': list(range(2, 100)), 'min_samples_split': [2, 3, 4]}
grid_search_cv = GridSearchCV(DecisionTreeClassifier(random_state=42), params,
grid_search_cv.fit(X_train, y_train)
```

Fitting 3 folds for each of 294 candidates, totalling 882 fits

```
Out[49]:
```


```
GridSearchCV
  estimator: DecisionTreeClassifier
    DecisionTreeClassifier
```

```
In [50]: y_pred_dec = grid_search_cv.predict(X_test)
```

```
In [51]: print(accuracy_score(y_test, y_pred_dec))
```

0.7931850449597728


```
In [52]: display = PrecisionRecallDisplay.from_estimator(  
    grid_search_cv, X_test, y_test, name="Dec_tree"  
)  
_ = display.ax_.set_title("2-class Precision-Recall curve")
```



```
In [53]: roc_auc_score(y_test, y_pred_dec)
```

Out[53]: 0.7015730378374406

```
In [54]: lr_fpr_dec, lr_tpr_dec, _ = roc_curve(y_test, y_pred_dec)
# plot the roc curve for the model
plt.plot(lr_fpr_dec, lr_tpr_dec, marker='.', label='Decision Tree Classifier')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
# show the legend
plt.legend()
# show the plot
plt.show()
```




```
In [55]: print(classification_report(y_test, y_pred_dec))
```

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1552
1	0.64	0.51	0.57	561
accuracy			0.79	2113
macro avg	0.74	0.70	0.71	2113
weighted avg	0.78	0.79	0.78	2113

```
In [56]: confusion_matrix_dec = confusion_matrix(y_test, y_pred_dec)

cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix_dec)

cm_display.plot()
plt.show()
```



Naive Bayes Classifier

```
In [57]: nb_classifier = GaussianNB()

params_NB = {'var_smoothing': np.logspace(0, -9, num=100)}
gs_NB = GridSearchCV(estimator=nb_classifier,
                      param_grid=params_NB,
                      cv=3,
                      verbose=1,
                      scoring='accuracy')
gs_NB.fit(X_train, y_train)

gs_NB.best_params_
```

Fitting 3 folds for each of 100 candidates, totalling 300 fits


```
Out[57]: {'var_smoothing': 1.0}
```

```
In [58]: y_pred_NB = gs_NB.predict(X_test)
```


```
In [59]: print(accuracy_score(y_test, y_pred_NB))
```

0.7406530998580217

```
In [60]: display = PrecisionRecallDisplay.from_estimator(  
        gs_NB, X_test, y_test, name="GaussianNB"  
    )  
_ = display.ax_.set_title("2-class Precision-Recall curve")
```



```
In [61]: lr_fpr_NB, lr_tpr_NB, _ = roc_curve(y_test, y_pred_NB)
# plot the roc curve for the model
plt.plot(lr_fpr_NB, lr_tpr_NB, marker='.', label='Decision Tree Classifier')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
# show the legend
plt.legend()
# show the plot
plt.show()
```




```
In [62]: print(classification_report(y_test, y_pred_NB))
```

	precision	recall	f1-score	support
0	0.90	0.73	0.80	1552
1	0.51	0.78	0.61	561
accuracy			0.74	2113
macro avg	0.70	0.75	0.71	2113
weighted avg	0.80	0.74	0.75	2113

```
In [63]: confusion_matrix_NB = confusion_matrix(y_test, y_pred_NB)

cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix_NB)

cm_display.plot()
plt.show()
```



Random Forest

```
In [64]: param_grid_rfc = {
    'n_estimators': [50, 150],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth' : [4,5,6,7,8],
    'criterion' :['gini', 'entropy']
}
```

```
In [65]: rfc=RandomForestClassifier(random_state=123)
```



```
In [66]: CV_rfc = GridSearchCV(estimator=rfc, param_grid = param_grid_rfc, cv= 3)
CV_rfc.fit(X_train, y_train)
```

E:\Anaconda Nav\lib\site-packages\sklearn\ensemble_forest.py:424: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'` or remove this parameter as it is also the default value for RandomForestClassifiers and ExtraTreesClassifiers.

warn(

E:\Anaconda Nav\lib\site-packages\sklearn\ensemble_forest.py:424: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'` or remove this parameter as it is also the default value for RandomForestClassifiers and ExtraTreesClassifiers.

warn(

E:\Anaconda Nav\lib\site-packages\sklearn\ensemble_forest.py:424: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'` or remove this parameter as it is also the default value for RandomForestClassifiers and ExtraTreesClassifiers.

warn(

E:\Anaconda Nav\lib\site-packages\sklearn\ensemble_forest.py:424: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'` or remove this parameter as it is also the default value for RandomForestClassifiers and ExtraTreesClassifiers.

```
In [67]: CV_rfc.best_params_
```

```
Out[67]: {'criterion': 'entropy',
          'max_depth': 8,
          'max_features': 'auto',
          'n_estimators': 150}
```

```
In [68]: rfc1=RandomForestClassifier(random_state=42, max_features='auto', n_estimators=
```

```
In [69]: rfc1.fit(X_train, y_train)
```

E:\Anaconda Nav\lib\site-packages\sklearn\ensemble_forest.py:424: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'` or remove this parameter as it is also the default value for RandomForestClassifiers and ExtraTreesClassifiers.

warn(


```
Out[69]: RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=8, max_features='auto',
                       n_estimators=200, random_state=42)
```

```
In [70]: y_pred_rfc = rfc1.predict(X_test)
```


```
In [71]: print(accuracy_score(y_test, y_pred_rfc))
```

```
0.8026502602934217
```

```
In [72]: display = PrecisionRecallDisplay.from_estimator(  
    rfc1, X_test, y_test, name="RandomForest"  
)  
_ = display.ax_.set_title("2-class Precision-Recall curve")
```



```
In [73]: lr_fpr_rfc, lr_tpr_rfc, _ = roc_curve(y_test, y_pred_rfc)
# plot the roc curve for the model
plt.plot(lr_fpr_rfc, lr_tpr_rfc, marker='.', label='Random Forest Classifier')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
# show the legend
plt.legend()
# show the plot
plt.show()
```




```
In [74]: print(classification_report(y_test, y_pred_rfc))
```

	precision	recall	f1-score	support
0	0.83	0.92	0.87	1552
1	0.68	0.49	0.57	561
accuracy			0.80	2113
macro avg	0.76	0.70	0.72	2113
weighted avg	0.79	0.80	0.79	2113

```
In [75]: confusion_matrix_rfc = confusion_matrix(y_test, y_pred_rfc)

cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix_rfc)

cm_display.plot()
plt.show()
```



Stochastic Gradient Classifier

```
In [76]: param_dist_SGD = {  
    'loss': ['hinge', 'log', 'perceptron'],  
    'alpha': [0.0001, 0.001, 0.01, 0.1],  
    'penalty': ['l2', 'l1', 'elasticnet']  
}  
  
model = SGDClassifier()  
randomized_search = RandomizedSearchCV(model, param_dist_SGD, n_iter=10, cv=5)  
randomized_search.fit(X_train, y_train)
```

[illegible]

[illegible]

Out[76]:

- ▶ **RandomizedSearchCV**
- ▶ **estimator: SGDClassifier**
 - ▶ SGDClassifier

```
In [77]: randomized_search.best_params_
```

```
Out[77]: {'penalty': 'l1', 'loss': 'log', 'alpha': 0.001}
```

```
In [78]: SGD1 = SGDClassifier(random_state = 42, penalty = "l2", loss = "log", alpha = 0.001)
SGD1.fit(X_train, y_train)
```

E:\Anaconda Nav\lib\site-packages\sklearn\linear_model_stochastic_gradient.py:163: FutureWarning: The loss 'log' was deprecated in v1.1 and will be removed in version 1.3. Use 'loss='log_loss'' which is equivalent.
warnings.warn(


```
Out[78]: SGDClassifier
SGDClassifier(alpha=0.001, loss='log', random_state=42)
```

```
In [79]: y_pred_SGD = SGD1.predict(X_test)
```

```
In [80]: print(accuracy_score(y_test, y_pred_SGD))
```

0.7927117841930904

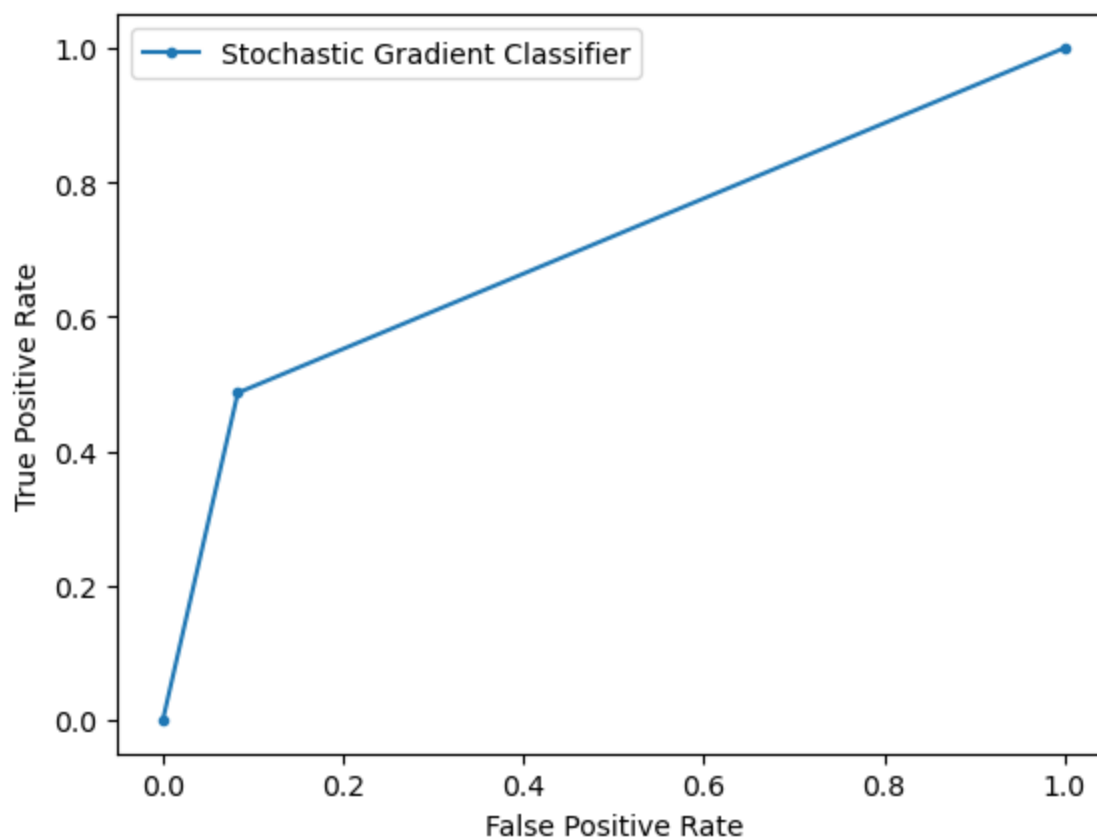
```
In [81]: display = PrecisionRecallDisplay.from_estimator(
    rfc1, X_test, y_test, name="StochasticGradient"
)
_ = display.ax_.set_title("2-class Precision-Recall curve")
```




```
In [82]: roc_auc_score(y_test, y_pred_SGD)
```

```
Out[82]: 0.7393806163515076
```

```
In [83]: lr_fpr_SGD, lr_tpr_SGD, _ = roc_curve(y_test, y_pred_rfc)
# plot the roc curve for the model
plt.plot(lr_fpr_SGD, lr_tpr_SGD, marker='.', label='Stochastic Gradient Classifier')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
# show the legend
plt.legend()
# show the plot
plt.show()
```




```
In [84]: print(classification_report(y_test, y_pred_SGD))
```

	precision	recall	f1-score	support
0	0.86	0.85	0.86	1552
1	0.61	0.63	0.62	561
accuracy			0.79	2113
macro avg	0.73	0.74	0.74	2113
weighted avg	0.79	0.79	0.79	2113

```
In [85]: confusion_matrix_SGD = confusion_matrix(y_test, y_pred_SGD)

cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix_SGD)

cm_display.plot()
plt.show()
```




Summary

```
In [87]: #calculate roc_auc score of each and save to a dataframe
roc_auc_scores = pd.DataFrame(columns=['Model', 'ROC_AUC'])
roc_auc_scores.loc[0] = ['svm_rbf', roc_auc_score(y_test, SVC_grid.predict(X_test))]
roc_auc_scores.loc[1] = ['tree_clf', roc_auc_score(y_test, grid_search_cv.predict(X_test))]
roc_auc_scores.loc[2] = ['randomforestclf', roc_auc_score(y_test, rfc1.predict(X_test))]
roc_auc_scores.loc[3] = ['bayes_clf', roc_auc_score(y_test, gs_NB.predict(X_test))]
roc_auc_scores.loc[4] = ['sgd_clf', roc_auc_score(y_test, SGD1.predict(X_test))]
roc_auc_scores
```

Out[87]:

	Model	ROC_AUC
0	svm_rbf	0.698685
1	tree_clf	0.701573
2	randomforestclf	0.701756
3	bayes_clf	0.752885
4	sgd_clf	0.739381

```
In [88]: #make a bar chart to show the highest values of roc_auc with values also printed
plt.figure(figsize=(10,5))
sns.barplot(x='Model',y='ROC_AUC',data=roc_auc_scores)
plt.xticks(rotation=90)
for i in range(len(roc_auc_scores)):
    plt.text(i,roc_auc_scores['ROC_AUC'][i],round(roc_auc_scores['ROC_AUC'][i],2))
plt.show()
```



Thus bayes_clf is the best estimator for Churn Modelling